Performance

Frontend Instance Class: F1 (600MHz, 128MB) ▾

Adjusting your Frontend Instance Class will affect all frontend versions of your application. Your frontends will have more memory and processing power, but also consume frontend instance hours at an increased rate, which will lead to increased costs. For example an F2 consumes instance hours at twice the rate of an F1. Learn more.

Idle Instances: (2 – 4)

The Idle Instances slider allows you to control the number of idle instances available to the default version of your application at any given time. Idle Instances are pre-loaded with your application code, so when a new Instance is needed, it can serve traffic immediately. You will not be charged for idle instances over the specified maximum. A smaller number of idle Instances means your application costs less to run, but may encounter more startup latency during load spikes. Learn more.

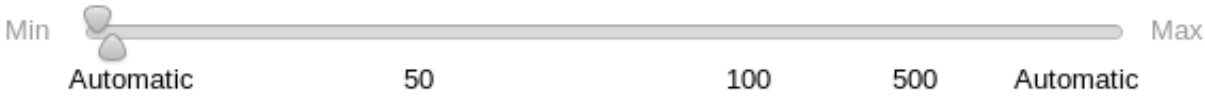

Pending Latency: (990ms – Automatic)

The Pending Latency slider controls how long requests spend in the pending queue before being served by an Instance of the default version of your application. If the minimum pending latency is high App Engine will allow requests to wait rather than start new Instances to process them. This can reduce the number of instance hours your application uses, but can result in more user-visible latency. Learn more.

PageSpeed Service:

The PageSpeed Service automatically optimizes and caches your site for improved performance. To configure advanced settings, edit the pagespeed section in your app.yaml file.

☐ Enable PageSpeed Service Flush Cache

Save Settings